

IMDb - Exploratory Data Analysis (SQL + Python)

The primary objective of this exploratory data analysis project is to gain insights into the IMDb dataset using SQL and Python. By leveraging the power of structured query language for efficient data manipulation and Python's versatile libraries for in-depth analysis and visualization, we aim to discover valuable information about the movies featured in the dataset. Through this comprehensive analysis, we seek to unravel the intricacies of the movie industry, from genre trends to the impact of directors, actors and many more.

Reading the dataset

```
In [231]: 1 # Importing all the necessary libraries
          2 import pandas as pd
          3 import numpy as np
          4 import matplotlib.pyplot as plt
          5 import seaborn as sns
          6 import sqlalchemy
          7 from matplotlib import style
```

```
In [232]: 1 # Connecting with PostgreSQL
          2 engine = sqlalchemy.create_engine('postgresql://postgres:Adi_1997@localhost:54
```

```
In [233]: 1 # Exploring the dataset
          2 df = pd.read_sql('IMDb',engine)
          3 df.head()
```

```
Out[233]:
```

	Title	Year	Certificate	Runtime	Genre	Rating	Meta_score	Director	Star1	Star2
0	The Shawshank Redemption	1994	A	142	Drama	9.3	80.0	Frank Darabont	Tim Robbins	Morgan Freeman
1	The Godfather	1972	A	175	Crime, Drama	9.2	100.0	Francis Ford Coppola	Marlon Brando	Al Pacino
2	The Dark Knight	2008	UA	152	Action, Crime, Drama	9.0	84.0	Christopher Nolan	Christian Bale	Heath Ledger
3	The Godfather: Part II	1974	A	202	Crime, Drama	9.0	90.0	Francis Ford Coppola	Al Pacino	Robert De Niro
4	12 Angry Men	1957	U	96	Crime, Drama	9.0	96.0	Sidney Lumet	Henry Fonda	Lee J. Cobb

```
In [234]: 1 df.columns
```

```
Out[234]: Index(['Title', 'Year', 'Certificate', 'Runtime', 'Genre', 'Rating',
                  'Meta_score', 'Director', 'Star1', 'Star2', 'Star3', 'Star4', 'Votes',
                  'Gross'],
                  dtype='object')
```

Movies

Total number of movies in the dataset

```
In [235]: 1 query = ''' SELECT COUNT(*) AS Total_no_of_movies
2             FROM "IMDb"
3             '''
```

```
In [236]: 1 df = pd.read_sql_query(query,engine)
2 df
```

```
Out[236]:
```

total_no_of_movies	
0	1000

Top 10 highest-rated movies

```
In [237]: 1 query = ''' SELECT "Title", "Rating"
2             FROM "IMDb"
3             ORDER BY "Rating" DESC
4             LIMIT 10'''
```

```
In [238]: 1 df = pd.read_sql_query(query,engine)
2 df
```

```
Out[238]:
```

	Title	Rating
0	The Shawshank Redemption	9.3
1	The Godfather	9.2
2	The Dark Knight	9.0
3	The Godfather: Part II	9.0
4	12 Angry Men	9.0
5	Pulp Fiction	8.9
6	The Lord of the Rings: The Return of the King	8.9
7	Schindler's List	8.9
8	Fight Club	8.8
9	Inception	8.8

Actors and Directors

Top 10 directors with the most movies in the dataset

```
In [239]: 1 query = ''' SELECT "Director" , COUNT(*) AS Total_movies
2           FROM "IMDb"
3           GROUP BY "Director"
4           ORDER BY Total_movies DESC
5           LIMIT 10
6           '''
```

```
In [240]: 1 df = pd.read_sql_query(query,engine)
2         df
```

```
Out[240]:
```

	Director	total_movies
0	Alfred Hitchcock	14
1	Steven Spielberg	13
2	Hayao Miyazaki	11
3	Martin Scorsese	10
4	Akira Kurosawa	10
5	Billy Wilder	9
6	Woody Allen	9
7	Stanley Kubrick	9
8	David Fincher	8
9	Clint Eastwood	8

Movies in which 'Al Pacino' has appeared

```
In [13]: 1 query = ''' SELECT "Title"
2           FROM "IMDb"
3           WHERE "Star1" = 'Al Pacino' OR "Star2" = 'Al Pacino' OR "Star3" =
4           '''
```

```
In [14]: 1 df = pd.read_sql_query(query,engine)
          2 df
```

```
Out[14]:
```

	Title
0	The Godfather
1	The Godfather: Part II
2	Scarface
3	Heat
4	Scent of a Woman
5	Dog Day Afternoon
6	The Irishman
7	Carlito's Way
8	The Insider
9	Donnie Brasco
10	Glengarry Glen Ross
11	Serpico
12	The Godfather: Part III

Lead actor who has worked with 'Chirstopher Nolan' the most

```
In [70]: 1 query = ''' SELECT "Star1" AS Actor, COUNT(*) AS Total_movies
          2           FROM "IMDb"
          3           WHERE "Director" = 'Christopher Nolan'
          4           GROUP BY Actor
          5           ORDER BY Total_movies DESC
          6           LIMIT 1
          7
          8           '''
```

```
In [71]: 1 df = pd.read_sql_query(query,engine)
          2 df
```

```
Out[71]:
```

	actor	total_movies
0	Christian Bale	4

Genres

Most popular genre in the dataset

```
In [241]: 1 query = ''' SELECT "Genre", COUNT (*) AS Total_movies
          2           FROM "IMDb"
          3           GROUP BY "Genre"
          4           ORDER BY Total_movies DESC
          5           LIMIT 15
          6           '''
```

In [242]:

1

df = pd.read_sql_query(query,engine)

2

df

Out[242]:


	Genre	total_movies
0	Drama	85
1	Drama, Romance	37
2	Comedy, Drama	35
3	Comedy, Drama, Romance	31
4	Action, Crime, Drama	30
5	Biography, Drama, History	28
6	Crime, Drama, Thriller	28
7	Crime, Drama, Mystery	27
8	Crime, Drama	26
9	Animation, Adventure, Comedy	24
10	Action, Adventure, Sci-Fi	21
11	Biography, Crime, Drama	16
12	Drama, War	15
13	Comedy, Crime, Drama	15
14	Action, Adventure, Drama	14

In [243]:

```
1 plt.figure(figsize= (14,6))
2 sns.barplot(x='Genre',y='total_movies', data=df, palette= 'deep')
3 plt.xticks(rotation = 90)
```

Out[243]:

```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]),
 [Text(0, 0, 'Drama'),
  Text(1, 0, 'Drama, Romance'),
  Text(2, 0, 'Comedy, Drama'),
  Text(3, 0, 'Comedy, Drama, Romance'),
  Text(4, 0, 'Action, Crime, Drama'),
  Text(5, 0, 'Biography, Drama, History'),
  Text(6, 0, 'Crime, Drama, Thriller'),
  Text(7, 0, 'Crime, Drama, Mystery'),
  Text(8, 0, 'Crime, Drama'),
  Text(9, 0, 'Animation, Adventure, Comedy'),
  Text(10, 0, 'Action, Adventure, Sci-Fi'),
  Text(11, 0, 'Biography, Crime, Drama'),
  Text(12, 0, 'Drama, War'),
  Text(13, 0, 'Comedy, Crime, Drama'),
  Text(14, 0, 'Action, Adventure, Drama')])
```



As we can see from the above barplot, Drama is the most popular genre in the dataset. Also, Drama is present in the movies which has multiple genres.

Top 5 genres with the highest average ratings (Min 10 movies)

In [336]:

```
1 query = ''' SELECT DISTINCT "Genre" , CAST(AVG("Rating") AS DECIMAL(10,2)) as
2             FROM "IMDb"
3             GROUP BY "Genre"
4             HAVING COUNT(*) >= 10
5             ORDER BY Avg_Rating DESC
6             LIMIT 5
7             '''
```

```
In [337]: 1 df = pd.read_sql_query(query, engine)
          2 df
```

```
Out[337]:
```

	Genre	avg_rating
0	Crime, Drama	8.16
1	Action, Adventure, Drama	8.15
2	Drama, War	8.07
3	Biography, Drama, History	8.02
4	Biography, Drama	7.98

Most profitable genres

```
In [333]: 1 query = ''' SELECT DISTINCT "Genre", CAST(AVG("Gross") AS BIGINT) AS Avg_Gross
          2           FROM "IMDb"
          3           WHERE "Gross" <> 0
          4           GROUP BY "Genre"
          5           ORDER BY Avg_Gross DESC
          6           LIMIT 5
          7           '''
```


```
In [334]: 1 df = pd.read_sql_query(query, engine)
          2 df
```

```
Out[334]:
```

	Genre	avg_gross
0	Family, Sci-Fi	435110554
1	Action, Adventure, Fantasy	352723505
2	Action, Adventure, Family	301959197
3	Action, Adventure, Sci-Fi	280888546
4	Adventure, Fantasy	280685212

```
In [335]: 1 plt.figure(figsize= (8,4))
2 sns.barplot(x='Genre',y='avg_gross', data=df, palette= 'gist_stern')
3 plt.xticks(rotation = 45)
```

```
Out[335]: (array([0, 1, 2, 3, 4]),
 [Text(0, 0, 'Family, Sci-Fi'),
  Text(1, 0, 'Action, Adventure, Fantasy'),
  Text(2, 0, 'Action, Adventure, Family'),
  Text(3, 0, 'Action, Adventure, Sci-Fi'),
  Text(4, 0, 'Adventure, Fantasy')])
```



Action and Adventure are two of the most profitable genres as per above bar plot.

Box Office Analysis

Top 10 highest-grossing movies

```
In [323]: 1 query = ''' SELECT "Title" ,"Director", "Gross", "Genre"
2 FROM "IMDb"
3 WHERE "Gross" <> 0
4 ORDER BY "Gross" DESC
5 LIMIT 10
6 '''
```


In [324]:

1 df = pd.read_sql_query(query,engine)


2 df

Out[324]:

	Title	Director	Gross	Genre
0	Star Wars: Episode VII - The Force Awakens	J.J. Abrams	936662225	Action, Adventure, Sci-Fi
1	Avengers: Endgame	Anthony Russo	858373000	Action, Adventure, Drama
2	Avatar	James Cameron	760507625	Action, Adventure, Fantasy
3	Avengers: Infinity War	Anthony Russo	678815482	Action, Adventure, Sci-Fi
4	Titanic	James Cameron	659325379	Drama, Romance
5	The Avengers	Joss Whedon	623279547	Action, Adventure, Sci-Fi
6	Incredibles 2	Brad Bird	608581744	Animation, Action, Adventure
7	The Dark Knight	Christopher Nolan	534858444	Action, Crime, Drama
8	Rogue One	Gareth Edwards	532177324	Action, Adventure, Sci-Fi
9	The Dark Knight Rises	Christopher Nolan	448139099	Action, Adventure

```
In [304]: 1 plt.figure(figsize= (14,8))
          2 sns.barplot(x='Title',y='Gross', data=df, palette= 'bright')
          3 plt.xticks(rotation = 90)
```

```
Out[304]: (array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]),
 [Text(0, 0, 'Star Wars: Episode VII - The Force Awakens'),
  Text(1, 0, 'Avengers: Endgame'),
  Text(2, 0, 'Avatar'),
  Text(3, 0, 'Avengers: Infinity War'),
  Text(4, 0, 'Titanic'),
  Text(5, 0, 'The Avengers'),
  Text(6, 0, 'Incredibles 2'),
  Text(7, 0, 'The Dark Knight'),
  Text(8, 0, 'Rogue One'),
  Text(9, 0, 'The Dark Knight Rises')])
```



From the above list of top grossers, we can see that there is only one movie which has a 'Drama' genre and rest of them are mostly of 'Action' or 'Adventure'. So, even if Drama is the most popular genre, Action and Adventure are two of the most profitable ones.

Director with the highest total box office revenue

```
In [305]: 1 query = ''' SELECT "Director" , CAST(SUM("Gross") AS BIGINT) AS Total_collecti
2          FROM "IMDb"
3          WHERE "Gross" <> 0
4          GROUP BY "Director"
5          ORDER BY Total_collection DESC
6          LIMIT 5
7          '''
```

```
In [306]: 1 df = pd.read_sql_query(query,engine)
2          df
```

```
Out[306]:
```

	Director	total_collection
0	Steven Spielberg	2478133165
1	Anthony Russo	2205039403
2	Christopher Nolan	1937454106
3	James Cameron	1748236602
4	Peter Jackson	1597312443

Box office revenue for each year

```
In [307]: 1 query = ''' SELECT "Year", CAST(SUM("Gross") AS BIGINT) AS Total_collections
2          FROM "IMDb"
3          WHERE "Gross" <> 0
4          GROUP BY "Year"
5          ORDER BY "Year" ASC
6          '''
```

```
In [308]: 1 df = pd.read_sql_query(query,engine)
2          df
```

```
Out[308]:
```

	Year	total_collections
0	1921	5450000
1	1924	977375
2	1925	5500970
3	1926	1033895
4	1927	1775706
...
89	2015	2462336868
90	2016	2595557425
91	2017	2061312852
92	2018	2607757362
93	2019	2406742688

94 rows × 2 columns


In [309]:

```
1 plt.figure(figsize= (14,8))
2 sns.lineplot(x='Year',y='total_collections', data=df, palette= 'inferno')
3 plt.xticks(rotation = 0)
```

C:\Users\Aditya\AppData\Local\Temp\ipykernel_6280\4263806761.py:2: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.

```
sns.lineplot(x='Year',y='total_collections', data=df, palette= 'inferno')
```

Out[309]: (array([1900., 1920., 1940., 1960., 1980., 2000., 2020., 2040.]),
[Text(1900.0, 0, '1900'),
Text(1920.0, 0, '1920'),
Text(1940.0, 0, '1940'),
Text(1960.0, 0, '1960'),
Text(1980.0, 0, '1980'),
Text(2000.0, 0, '2000'),
Text(2020.0, 0, '2020'),
Text(2040.0, 0, '2040')])



From the above time series analysis of box office revenue, we can say that last two decades have seen significant growth in the box office collection.

User Reviews

Top 10 movies with the most user reviews

In [255]:

```
1 query = ''' SELECT "Title", "Votes", "Rating"
2             FROM "IMDb"
3             ORDER BY "Votes" DESC
4             LIMIT 10
5             '''
```

```
In [256]: 1 df = pd.read_sql_query(query,engine)
          2 df
```

```
Out[256]:
```

	Title	Votes	Rating
0	The Shawshank Redemption	2343110	9.3
1	The Dark Knight	2303232	9.0
2	Inception	2067042	8.8
3	Fight Club	1854740	8.8
4	Pulp Fiction	1826188	8.9
5	Forrest Gump	1809221	8.8
6	The Matrix	1676426	8.7
7	The Lord of the Rings: The Fellowship of the Ring	1661481	8.8
8	The Lord of the Rings: The Return of the King	1642758	8.9
9	The Godfather	1620367	9.2

Movies with the highest and lowest user ratings

```
In [257]: 1 query = ''' WITH cte AS(
          2 SELECT "Title",
          3 "Rating" AS Highest,
          4 ROW_NUMBER () OVER (ORDER BY "Rating" DESC ) AS Rank_high
          5 FROM "IMDb"
          6 GROUP BY "Title", "Rating"
          7 ),
          8
          9 cte1 AS(
         10 SELECT "Title",
         11 "Rating" AS Lowest,
         12 ROW_NUMBER () OVER (ORDER BY "Rating" ASC ) AS Rank_low
         13 FROM "IMDb"
         14 GROUP BY "Title", "Rating"
         15 )
         16
         17 SELECT cte."Title", cte.Highest, cte1.Lowest
         18 FROM cte JOIN cte1 ON cte."Title" = cte1."Title"
         19 WHERE cte.Rank_high = 1 OR cte1.Rank_low = 1
         20 ORDER BY "Title" DESC
         21
         22 '''
```

```
In [258]: 1 df = pd.read_sql_query(query,engine)
          2 df
```

```
Out[258]:
```

	Title	highest	lowest
0	The Shawshank Redemption	9.3	9.3
1	The Secret of Kells	7.6	7.6

Correlation between user ratings and box office revenue

```
In [326]: 1 query = ''' SELECT "Rating", "Gross"
           2           FROM "IMDb"
           3           WHERE "Gross" <> 0
           4           '''
```

```
In [327]: 1 df = pd.read_sql_query(query,engine)
           2 df
```

```
Out[327]:
```

	Rating	Gross
0	9.3	28341469
1	9.2	134966411
2	9.0	534858444
3	9.0	57300000
4	9.0	4360000
...
826	7.6	696690
827	7.6	1378435
828	7.6	141843612
829	7.6	13780024
830	7.6	30500000

831 rows × 2 columns

```
In [330]: 1 corr = df["Rating"].corr(df["Gross"])
           2 corr
```


```
Out[330]: 0.09592277110132366
```

In [331]:

```
1 plt.figure(figsize= (10,6))
2 sns.scatterplot(x='Rating',y='Gross', data=df)
3 plt.xticks(rotation = 0)
```

Out[331]:

```
(array([7.5 , 7.75, 8. , 8.25, 8.5 , 8.75, 9. , 9.25, 9.5 ]),
 [Text(7.5, 0, '7.50'),
  Text(7.75, 0, '7.75'),
  Text(8.0, 0, '8.00'),
  Text(8.25, 0, '8.25'),
  Text(8.5, 0, '8.50'),
  Text(8.75, 0, '8.75'),
  Text(9.0, 0, '9.00'),
  Text(9.25, 0, '9.25'),
  Text(9.5, 0, '9.50')])
```



Correlation coefficient between user ratings and box office revenue is 0.096, which is close to 0, so we can say that user rating has no direct impact on box office collection and vice versa.

Director and Cast Analysis

Average rating for each director's movies (Min 5 movies)

In [338]:

```
1 query = ''' SELECT "Director", CAST(AVG("Rating") AS DECIMAL(10,2)) AS Avg_Rat
2             FROM "IMDb"
3             GROUP BY "Director"
4             HAVING COUNT(*) >= 5
5             ORDER BY Avg_Rating DESC
6             '''
```

In [339]:

```
1 df = pd.read_sql_query(query,engine)
2 df
```

Out[339]:

	Director	avg_rating
0	Christopher Nolan	8.46
1	Peter Jackson	8.40
2	Francis Ford Coppola	8.40
3	Charles Chaplin	8.33
4	Sergio Leone	8.27
5	Stanley Kubrick	8.23
6	Akira Kurosawa	8.22
7	Quentin Tarantino	8.18
8	Martin Scorsese	8.17
9	Billy Wilder	8.14
10	Ingmar Bergman	8.14
11	Andrei Tarkovsky	8.12
12	Robert Zemeckis	8.12
13	Sidney Lumet	8.10
14	James Cameron	8.08
15	Ridley Scott	8.08
16	David Fincher	8.04
17	Steven Spielberg	8.03
18	Hayao Miyazaki	8.02
19	Alfred Hitchcock	8.01
20	Federico Fellini	8.00
21	Roman Polanski	8.00
22	Denis Villeneuve	7.98
23	Ron Howard	7.92
24	Clint Eastwood	7.91
25	Richard Linklater	7.90
26	John Ford	7.90
27	John Huston	7.90
28	Howard Hawks	7.86
29	David Lynch	7.86
30	Rob Reiner	7.83
31	Wes Anderson	7.83
32	Joel Coen	7.82
33	Woody Allen	7.79
34	Alfonso Cuarón	7.75

Christopher Nolan's movies as a director has the highest average rating for directors who have directed minimum 5 movies.

Lead actors who are associated with higher-grossing movies

```
In [315]: 1 query = ''' SELECT "Star1" AS Actor, CAST(SUM("Gross") AS BIGINT) AS total_box
2          2 FROM "IMDb"
3          3 WHERE "Gross" <> 0
4          4 GROUP BY "Star1"
5          5 ORDER BY total_boxoffice DESC
6          6 LIMIT 10
7          7 '''
```

```
In [316]: 1 df = pd.read_sql_query(query, engine)
2          2 df
```

```
Out[316]:
```


	actor	total_boxoffice
0	Tom Hanks	2493097454
1	Joe Russo	2205039403
2	Leonardo DiCaprio	1877321752
3	Daniel Radcliffe	1835901034
4	Christian Bale	1351591432
5	Robert Downey Jr.	1150720327
6	Elijah Wood	1035942020
7	Daisy Ridley	936662225
8	Mark Hamill	922340616
9	Craig T. Nelson	870022836

In [317]:

```
1 plt.figure(figsize= (14,8))
2 sns.barplot(x='actor',y='total_boxoffice', data=df, palette= 'tab10')
3 plt.xticks(rotation = 45)
```

Out[317]:

```
(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]),
 [Text(0, 0, 'Tom Hanks'),
  Text(1, 0, 'Joe Russo'),
  Text(2, 0, 'Leonardo DiCaprio'),
  Text(3, 0, 'Daniel Radcliffe'),
  Text(4, 0, 'Christian Bale'),
  Text(5, 0, 'Robert Downey Jr.'),
  Text(6, 0, 'Elijah Wood'),
  Text(7, 0, 'Daisy Ridley'),
  Text(8, 0, 'Mark Hamill'),
  Text(9, 0, 'Craig T. Nelson')])
```



Tom Hanks, Joe Russo and Leonardo DiCaprio are few lead actors who have been associated with high grossing movies.

Runtime Analysis

Average runtime of movies over the years

In [318]:

```
1 query = ''' SELECT "Year", CAST(AVG("Runtime") AS DECIMAL(10,2)) AS Avg_Runtim
2           FROM "IMDb"
3           GROUP BY "Year"
4           ORDER BY "Year"
5           '''
```

In [319]:

1 df = pd.read_sql_query(query, engine)

2 df

Out[319]:

	Year	avg_runtime
0	1920	76.00
1	1921	68.00
2	1922	94.00
3	1924	45.00
4	1925	85.00
...
94	2016	123.64
95	2017	121.59
96	2018	128.11
97	2019	132.13
98	2020	126.67

99 rows × 2 columns


In [320]:

```
1 plt.figure(figsize= (14,8))
2 sns.lineplot(x='Year',y='avg_runtime', data=df, palette= 'inferno')
3 plt.xticks(rotation = 0)
```

C:\Users\Aditya\AppData\Local\Temp\ipykernel_6280\2765243330.py:2: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.

```
sns.lineplot(x='Year',y='avg_runtime', data=df, palette= 'inferno')
```

Out[320]: (array([1900., 1920., 1940., 1960., 1980., 2000., 2020., 2040.]),
[Text(1900.0, 0, '1900'),
Text(1920.0, 0, '1920'),
Text(1940.0, 0, '1940'),
Text(1960.0, 0, '1960'),
Text(1980.0, 0, '1980'),
Text(2000.0, 0, '2000'),
Text(2020.0, 0, '2020'),
Text(2040.0, 0, '2040')])



From the above time series analysis of average runtime of movies, we can see that over the years there are lot of fluctuations in the average runtime from 45 mins to 150 mins but from the past 2 decades average runtime has consolidated around 120 mins.

Movies with the Longest and Shortest runtimes

```
In [270]: 1 query = ''' WITH cte AS(
2 SELECT "Title",
3 "Runtime" AS Longest,
4 ROW_NUMBER () OVER (ORDER BY "Runtime" DESC ) AS Rank_high
5 FROM "IMDb"
6 GROUP BY "Title", "Runtime"
7 ),
8
9 cte1 AS(
10 SELECT "Title",
11 "Runtime" AS Shortest,
12 ROW_NUMBER () OVER (ORDER BY "Runtime" ASC ) AS Rank_low
13 FROM "IMDb"
14 GROUP BY "Title", "Runtime"
15 )
16
17 SELECT cte."Title", cte.Longest, cte1.Shortest
18 FROM cte JOIN cte1 ON cte."Title" = cte1."Title"
19 WHERE cte.Rank_high = 1 OR cte1.Rank_low = 1
20 ORDER BY "Title" ASC '''
```

```
In [271]: 1 df = pd.read_sql_query(query, engine)
2 df
```

```
Out[271]:
```

	Title	longest	shortest
0	Gangs of Wasseypur	321	321
1	Sherlock Jr.	45	45

Movies with a runtime longer than the average duration

```
In [272]: 1 query = ''' SELECT "Title", "Runtime"
2 FROM "IMDb"
3 WHERE "Runtime" > (SELECT AVG("Runtime") FROM "IMDb")
4 ORDER BY "Runtime" DESC
5 '''
```

```
In [273]: 1 df = pd.read_sql_query(query, engine)
          2 df
```

```
Out[273]:
```

	Title	Runtime
0	Gangs of Wasseypur	321
1	Hamlet	242
2	Gone with the Wind	238
3	Once Upon a Time in America	229
4	Lawrence of Arabia	228
...
437	About Time	123
438	Jodaeiye Nader az Simin	123
439	The Theory of Everything	123
440	Atonement	123
441	The Notebook	123

442 rows × 2 columns

IMDb rating Vs Meta score

```
In [274]: 1 query = ''' SELECT "Rating", "Meta_score"
          2           FROM "IMDb"
          3           WHERE "Meta_score" <> 0 AND "Rating" <> 0
          4           '''
```

```
In [275]: 1 df = pd.read_sql_query(query, engine)
          2 df
```

```
Out[275]:
```

	Rating	Meta_score
0	9.3	80
1	9.2	100
2	9.0	84
3	9.0	90
4	9.0	96
...
838	7.6	76
839	7.6	84
840	7.6	85
841	7.6	78
842	7.6	93


843 rows × 2 columns

```
In [276]: 1 corr = df["Rating"].corr(df["Meta_score"])
          2 corr
```

```
Out[276]: 0.26853084455955467
```

```
In [277]: 1 plt.figure(figsize= (10,6))
          2 sns.scatterplot(x='Meta_score',y='Rating', data=df)
          3 plt.xticks(rotation = 0)
```

```
Out[277]: (array([ 20.,  30.,  40.,  50.,  60.,  70.,  80.,  90., 100., 110.]),
 [Text(20.0, 0, '20'),
  Text(30.0, 0, '30'),
  Text(40.0, 0, '40'),
  Text(50.0, 0, '50'),
  Text(60.0, 0, '60'),
  Text(70.0, 0, '70'),
  Text(80.0, 0, '80'),
  Text(90.0, 0, '90'),
  Text(100.0, 0, '100'),
  Text(110.0, 0, '110')])
```



Correlation coefficient between user ratings and Meta score is 0.268, which is closer to 0 than 1, so we can say that user rating has weak correlation with Meta score and vice versa.

Conclusion

In conclusion, our exploratory data analysis project on the IMDb movie dataset has unveiled a plethora of intriguing insights into the world of cinema. Notably, we've discovered that while drama stands as the most popular genre among the movies in our dataset, the genres that reign supreme in terms of profitability are action and adventure. This underscores the nuanced relationship between popularity and financial success in the film industry.

Our analysis of box office revenue over time has revealed a remarkable growth trend in the past two decades, reflecting the dynamic nature of the movie business. However, the correlation coefficient between user ratings and box office revenue suggests a weak connection, indicating that audience

preferences do not directly dictate a movie's financial performance.

Further, we found that Christopher Nolan's directorial ventures consistently earn the highest average ratings, emphasizing his influence in the industry. Additionally, actors like Tom Hanks, Joe Russo, and Leonardo DiCaprio have been associated with high-grossing movies, highlighting the impact of star power on a film's earnings.

The analysis of average movie runtime showed fluctuations over the years, with a consolidation around 120 minutes in the past two decades. Finally, the correlation between user ratings and Meta scores is weak, implying that audience opinions and critical acclaim often diverge.

Our EDA project has not only provided valuable insights into the movie industry's dynamics but also challenged preconceived notions, showcasing the complex interplay of factors that contribute to a movie's success. As the cinematic landscape continues to evolve, these findings serve as a testament to the ever changing nature of the art of filmmaking.

In []:	1	
---------	---	--